

✓ **Congratulations! You passed!**

Grade
received **80%**

Latest Submission
Grade 80%

To pass 80% or
higher

Go to next item

1. Face verification requires comparing a new picture against one person's face, whereas face recognition requires comparing a new picture against K persons' faces.

1 / 1 point

☒ True

☐ False

↩ Expand

✓ **Correct**
Correct.

2. Why is the face verification problem considered a one-shot learning problem? Choose the best answer.

1 / 1 point

- ☐ Because we are trying to compare to one specific person only.
- ☐ Because we have only have to forward pass the image one time through our neural network for verification.
- ☐ Because of the sensitive nature of the problem, we won't have a chance to correct it if the network makes a mistake.
- ☒ Because we might have only one example of the person we want to verify.

↩ Expand

✓ **Correct**
Correct. One-shot learning refers to the amount of data we have to solve a task.

3. You want to build a system that receives a person's face picture and determines if the person is inside a workgroup. You have pictures of all the faces of the people currently in the workgroup, but some members might leave, and some new members might be added. To train a system to solve this problem using the triplet loss you get many persons and take several pictures of each one. Which of the following do you agree with? (Select the best answer.)

0 / 1 point

- ☐ You shouldn't use persons outside the workgroup you are interested in because that might create a high variance in your model.
- ☐ You take several pictures of the same person to train $d(\text{img}_1, \text{img}_2)$ using the triplet loss.
- ☐ It would be best to increase the number of persons in the dataset by taking only one picture of each person to have a more representative set of the population.
- ☒ You take several pictures of the same person because this way you can get more pictures to train the network efficiently since you already have the person in place.

↩ Expand

✗ **Incorrect**
To train using the triplet loss you need several pictures of the same person, so you don't do this only to increase the size of the dataset.

4. Which of the following is a correct definition of the triplet loss? Consider that $\alpha > 0$. (We encourage you to figure out the answer from first principles, rather than just refer to the lecture.)

1 / 1 point

- ☐ $\max(\|f(A) - f(N)\|^2 - \|f(A) - f(P)\|^2 - \alpha, 0)$
- ☒ $\max(\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha, 0)$


- ☐ $\max(\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 - \alpha, 0)$
- ☐ $\max(\|f(A) - f(N)\|^2 - \|f(A) - f(P)\|^2 + \alpha, 0)$

Expand

✓ Correct
Correct

5. Consider the following Siamese network architecture:

1 / 1 point



The upper and lower neural networks have different input images, but have exactly the same parameters.

- ☐ False
- ☒ True

Expand

✓ Correct
Yes it is true, parameters are shared among these two networks.

6. You train a ConvNet on a dataset with 100 different classes. You wonder if you can find a hidden unit which responds strongly to pictures of cats. (I.e., a neuron so that, of all the input/training images that strongly activate that neuron, the majority are cat pictures.) You are more likely to find this unit in layer 4 of the network than in layer 1.

1 / 1 point

- ☒ True
- ☐ False

Expand

✓ Correct
Yes, this neuron understands complex shapes (cat pictures) so it is more likely to be in a deeper layer than in the first layer.

7. In neural style transfer, we train the pixels of an image, and not the parameters of a network.

1 / 1 point

- ☐ False
- ☒ True

Expand

✓ Correct
Correct. Neural style transfer compares the high-level features of two images and modifies the pixels of one of them in order to look artistic.

8. In neural style transfer the content loss J_{cont} is computed as:

0 / 1 point

$$J_{cont}(G, C) = \|a^{[l](C)} - a^{[l](G)}\|^2$$

Where $a^{[l](k)}$ is the activation of the l -th layer of a ConvNet trained for classification. We choose l to be a very high value to use compared to the more abstract activation of each image. True/False?

- ☐ False
- ☒ True

[Expand](#)

✗ **Incorrect**

We don't use a very deep layer since this will only compare if the two images belong to the same category.

9. In neural style transfer, what is updated in each iteration of the optimization algorithm?

1 / 1 point

- ☒ The pixel values of the generated image G
- ☐ The neural network parameters
- ☐ The pixel values of the content image C
- ☐ The regularization parameters

[Expand](#)

✓ **Correct**

Yes, neural style transfer is different from many of the algorithms you've seen up to now, because it doesn't learn any parameters; instead it learns directly the pixels of an image.

10. You are working with 3D data. You are building a network layer whose input volume has size $32 \times 32 \times 32 \times 16$ (this volume has 16 channels), and applies convolutions with 32 filters of dimension $3 \times 3 \times 3 \times 16$ (no padding, stride 1). What is the resulting output volume?

1 / 1 point

- ☐ Undefined: This convolution step is impossible and cannot be performed because the dimensions specified don't match up.
- ☐ $30 \times 30 \times 30 \times 16$
- ☒ $30 \times 30 \times 30 \times 32$

[Expand](#)

✓ **Correct**

Correct, you have used the formula $\lfloor \frac{n^{[l-1]} - f + 2 \times p}{s} \rfloor + 1 = n^{[l]}$ over the three first dimensions of the input data.